TriageNorth problem statement

Infrastructure teams receive symptoms from many monitoring systems, but resolving an incident still requires a person to assemble the surrounding facts. A container restarts, a host stops responding, and an application raises errors. These may describe one failure, yet arrive as separate alerts with different names and identifiers. Meanwhile, containers move, addresses change, and the failed instance may disappear before anyone investigates.

The operator must identify the affected resource, distinguish repeat notifications from new problems, reconstruct dependencies, retrieve logs and system state, and decide who may act. Evidence, proposed commands, approvals, and the result often live in separate tools. External ticket systems can drift from the incident, and opaque filtering makes it hard to see why a signal was discarded or grouped. Repeating that work slows response and makes automation difficult to trust.

TriageNorth turns infrastructure signals into a durable incident with an identifiable asset, an explanation supported by captured evidence, and a controlled next action. It gives operators a consistent workspace across a data center, cloud resources, hosts, and containers. Every change remains attributable to an authorized person or explicitly delegated automation principal.

The hackathon will demonstrate the complete path from a controlled failure to verified recovery. Success is a reproducible incident record and execution receipt, with measurable noise reduction and visible user control. No performance savings are claimed before measurement.